

§2 Background & preliminaries

- 2.1 Number formats & effective bits
- 2.2 Binary neural-network lineage
- 2.3 Straight-through estimator
- 2.4 Transformer-specific obstacles
- 2.5 Scaling behavior
- 2.6 Evaluation metrics

§3 Scope & corpus

- 3.1 Scope & eligibility
- 3.2 How the corpus was assembled
- 3.3 Effective bits & evidence basis
- 3.4 Limitations

§4 Taxonomy of binary & ternary LLMs

- 4.1 Weight representation
- 4.2 Training paradigm
- 4.3 Quantized components
- 4.4 Optimization mechanism
- 4.5 Systems stack
- 4.6 Secondary cut by modality
- 4.7 Reading the taxonomy

§5 Architectural design

- 5.1 The BitLinear primitive
- 5.2 Ternary versus binary in practice
- 5.3 Activation quantization
- 5.4 KV cache
- 5.5 Beyond the dense decoder
- 5.6 Matmul-free designs

§6 Training methodology & theory

- 6.1 The native training loop
- 6.2 Recipes
- 6.3 The STE and its alternatives
- 6.4 Distillation
- 6.5 Post-training routes to one bit
- 6.6 Scaling laws at one bit
- 6.7 Stability & smoothness

§7 Inference systems & hardware

- 7.1 The need for a software stack
- 7.2 bitnet.cpp & CPU kernels
- 7.3 GPU
- 7.4 Custom accelerators
- 7.5 Summary of the systems corpus

§8 Performance & benchmarking

- 8.1 The comparison problem
- 8.2 Perplexity & the parity claim
- 8.3 Downstream benchmarks
- 8.4 Reasoning & generation quality
- 8.5 Efficiency
- 8.6 A common-axis comparison
- 8.7 Summary of the evidence

§9 Applications & extensions

- 9.1 On-device & edge deployment
- 9.2 Embedding models
- 9.3 Speech
- 9.4 Vision-language & VLA models
- 9.5 Mixture-of-experts
- 9.6 Ecosystem beyond one laboratory
- 9.7 Environmental framing

§10 Independent re-evaluation

- 10.1 Rationale & scope
- 10.2 Setup
- 10.3 Accuracy results
- 10.4 Efficiency results
- 10.5 Summary of findings
- 10.6 Reproducibility audit
- 10.7 Threats

§11 Challenges & open problems

- 11.1 Capacity ceiling & token budget
- 11.2 Training stability at scale
- 11.3 The activation-precision floor
- 11.4 A theory of STE training
- 11.5 Hardware co-design
- 11.6 Reasoning & smoothness
- 11.7 Robustness, safety, security
- 11.8 Sub-1.58-bit representations
- 11.9 Evaluation & contamination

§12 Conclusion